

PRODYNA



Terraform



> Agenda

- Basic
 - What is Terraform?
 - Providers
 - Build, Update, Destroy
 - Input & Output Variables
 - Modules
- Advanced
 - Work in a team
 - State manipulation
 - Import Resources
 - Misc.

➤ Introduction



Dennis Creutz
Senior IT Consultant / Software Architect





- Name
- What you do?
- Experience with CI/CD/Terraform
- Expectations

OVERVIEW

TERRAFORM

➤ Infrastructure as Code

Infrastructure as code is the process of managing and provisioning computer data centers through machine-readable definition files

➤ Infrastructure as Code - Advantages

- Enables automation
 - Enables CI/CD of infrastructure
- Cost reduction
- Increased speed
- Risk reduction

- Chef
- Puppet
- Ansible
- CloudFormation, Azure Resource Manager, etc...
- Terraform

- Open-source infrastructure as code software
- Created by HashiCorp
- Written in HCL or JSON
- Supports many providers (AWS, GCP, Azure, Kubernetes, etc.)
- Current version: 0.12

➤ Install Terraform

- MacOS:
brew install terraform
- Windows:
<https://www.terraform.io/downloads.html>
- Verify installation:
terraform --version

BUILD, UPDATE AND DESTROY

TERRAFORM

➤ Terraform Provider

- Responsible for creating and managing resources
- Translator from HCL/JSON to API interactions
- Multiple providers in one Terraform file possible

> Terraform Provider

- AWS
- Azure
- GCP
- Kubernetes
- Helm
- MySQL
- Grafana
- CloudFlare
- Many many more: <https://www.terraform.io/docs/providers/index.html>

➤ Terraform Resource

Resource that exists within the infrastructure (e.g. EC2 instance)

```
resource "aws_vpc" "example" {  
    cidr_block = "10.0.0.0/16"  
}
```

„aws_vpc“ = Resource type, defined by the provider

„example“ = Resource name, defined by you

- Terraform will read all *.tf in a directory
 - Best practice: Always start with a main.tf
- `terraform init`
 - Initializes various local settings and data needed by other commands
 - Especially: Downloads all provider binaries
- `terraform plan [-out=myplan.tfplan]`
 - Creates a plan to visualise changes to the current infrastructure
 - No changes are applied!
 - Can be used as basic for the `terraform apply` command

- **terraform apply**
 - Creates a plan to visualise changes to the current infrastructure
 - Applies changes to your infrastructure (after you confirmed the changes)
- **terraform destroy**
 - Destroys all resources managed by Terraform
- **terraform fmt**
 - Formats all *.tf files
- **terraform validate**
 - Checks if the configuration is syntactically valid

DEMO & PRACTICE: CREATE

TERRAFORM

- macOS:
 - brew install awscli
- Windows:
 - https://docs.aws.amazon.com/de_de/cli/latest/userguide/install-windows.html
- Verify:
 - aws --version

Use AWS credentials in Terraform:

1. Login into AWS console
2. Create a user “terraform” with programmatic access and administrator policy
3. Download credentials
4. Configure AWS CLI with „aws configure“
 - Use your credentials
 - Region: eu-central-1
 - Output format: json

5. Create directory “my-tf-first-steps”
6. Create a file “aws_credentials.auto.tfvars”:

```
aws_credentials = {  
    access_key = "XXXXXX"  
    secret_key = "YYYYYYY"  
}
```

7. Create file “main.tf”
8. Add credential variable (more to variables later):

```
variable "aws_credentials" {  
    type = object({  
        access_key = string  
        secret_key = string  
    })  
    description = "AWS credentials used for terraform."  
}
```

9. Add AWS provider:

```
provider "aws" {  
    access_key = var.aws_credentials.access_key  
    secret_key = var.aws_credentials.secret_key  
    region = "eu-central-1"  
}
```

In „main.tf“:

1. Add a VPC:

```
resource "aws_vpc" "example" {  
    cidr_block = "10.0.0.0/16"  
}
```

2. terraform init
3. terraform apply

➤ Terraform State

- `terraform.tfstate`
- Maps real world resources to your configuration
- Delete state = Resources still exists but not managed by Terraform anymore
- Never interact with the state file directly!

DEMO & PRACTICE: UPDATE

TERRAFORM

› Demo: Update

1. Display your Terraform state:

`terraform show`

2. Change the CIDR of your VPC to 10.0.0.0/20

3. Create a Terraform plan:

`terraform plan -out=myExample.tfplan`

4. Apply plan:

`terraform apply myExample.tfplan`

DEMO & PRACTICE: DESTROY

TERRAFORM

› Demo: Destroy

1. Destroy everything
terraform destroy
2. Display Terraform State
terraform show

INPUT, OUTPUT & DEPENDENCIES

TERRAFORM

Defines Variables:

```
variable "myVar" {  
    default = "test"  
    type = string  
    description = "My example"  
}
```

Type can be string, number, object, list, bool and more.
No default value = required variable

➤ Terraform Variables

- Use Variables with „var.myVar“

- In Strings with \${var.myVar}

- Mostly used in Modules (later)

- Define local values:

```
locals {  
    stageName = "prod"  
    projectName = "prodyna-aws-training"  
}
```

Reference with „local.myLocal“

- Best practice: Define all variables in extra file „variables.tf“

Assignee Variables:

1. Command-line flags:

```
terraform apply -var 'myVar=test123'
```

2. From a file (like the AWS credentials)

- Files named *.auto.tfvars or terraform.tfvars are automatically propagated

3. From environment variables (only Strings):

```
TF_VAR_myVar=test123
```

➤ Terraform File Structure

- main.tf
 - Provider
 - Resources
- variables.tf
 - Only for variables and locals

```
output "database" {  
    value = aws_db_instance.this  
    sensitive = true  
    depends_on = [aws_db_instance.this]  
}
```

- Define outputs to..
 - Expose resource values (Modules..again)
 - Print values in terminal (terraform plan & apply)
- sensitive
If true, then no output in terminal
- Attention: All resource attributes and outputs are saved in plain text in the terraform state!

➤ Terraform File Structure

- main.tf
 - Provider
 - Resources
- variables.tf
 - Only for variables and locals
- output.tf
 - Only for outputs

➤ Terraform Remote State & Data Source

You can access other states (read only) via data source:

```
data "terraform_remote_state" "vpc" {  
    backend = "local"  
  
    config = {  
        path = "${path.module}/path/to/vpc/state/terraform.tfstate"  
    }  
}
```

- Data sources are only to read data and not to create resources
- Access output of the remote state:
`data.terraform_remote_state.vpc.outputs.myOutputName`
- Visit the provider docs to list all available data sources

➤ Terraform Dependencies

- Implicit:
 - Resource uses a value of another resource
- Explicit
 - Resource defines dependencies over „depends_on“

„The count parameter on resources can simplify configurations and let you scale resources by simply incrementing a number“

- Use to create multiple resources:

```
resource "aws_db_instance" "this" {  
    count = 3  
    name = "db-instance-${count.index}"  
}
```

- Or to add a condition to the creation of the resource:

```
resource "aws_vpc" "example" {  
    count = local.create_vpc ? 1 : 0  
  
    ...  
}
```

DEMO & PRACTICE: INPUT, OUTPUT & DEPENDENCIES

TERRAFORM

1. Extract the variables from “main.tf” into a new file “variables.tf”
2. Create a file “outputs.tf” and output the VPC values:

```
output "vpc" {  
    value = aws_vpc.example  
}
```

3. terraform apply
4. terraform show

5. Create multiple VPC's:

```
resource "aws_vpc" "example" {  
    count = 3  
  
    cidr_block = "10.${count.index}.0.0/20"  
}
```

6. What will happen?

7. terraform apply

PRACTICE: BASICS I

TERRAFORM

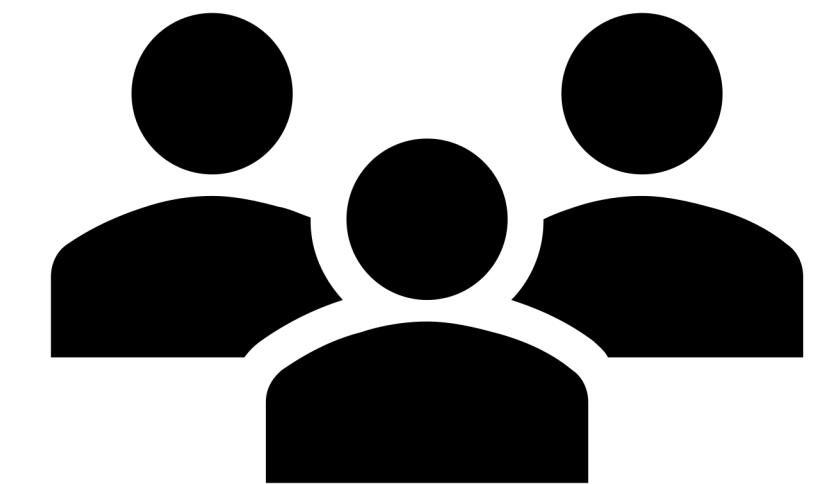
➤ Practice 1: Convert the shopping system to Terraform

Git repro: <https://github.com/DennisCreutz/prodyna-aws-training>

1. Use the code under “prodyna-aws-training/practices/terraform-basic-practice/” as basis
2. Copy the unfinished files and complete them.
3. Check for TODO's and missing resources.

The recommended order is:

1. VPC
 2. Database
 3. Backend
 4. API
 5. Frontend
-
4. Use the Terraform AWS provider docs: <https://www.terraform.io/docs/providers/aws/index.html>



Time: 120 min

MODULES

TERRAFORM

“A module is a container for multiple resources that are used together.”

- We already used a module, the root module
- Modules can use other modules:

```
module "vm" {  
    source = "./ec2-vm-module"  
    instances = 2  
}
```

- „source“
 - Required
 - Local or remote (like Git) location
- Other parameters are defined by the modules input variables
- You can access all modules output variables: `module.moduleName.varName`
- No built in „depends_on“ for Modules (yet)!

DEMO: MODULES

TERRAFORM

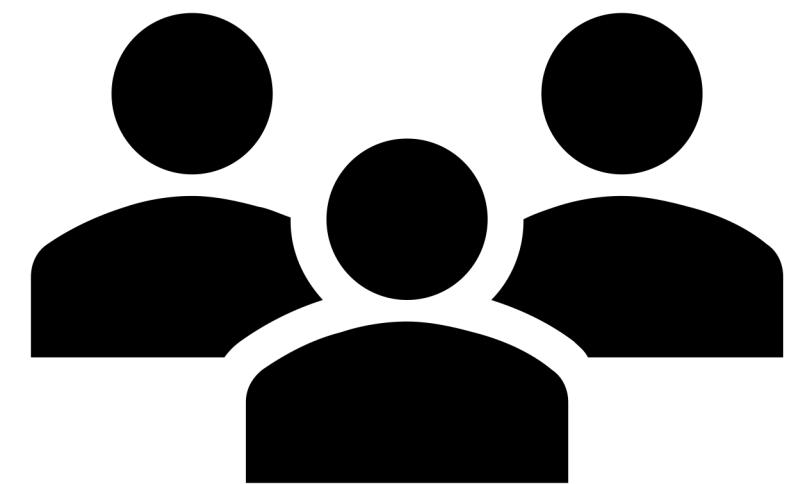
PRACTICE: BASICS II

TERRAFORM

➤ Practice 2: Modules

Extract the VPC as module

- Define all needed (input) variables and outputs.
- Reference your module



Time: 45 min

STATE MANIPULATION

TERRAFORM

➤ Terraform State Manipulation

- Output the Terraform state:
`terraform show`
- Recreate a already created resource:
`terraform taint addressOfTheResource`
- Apply/Destroy only specific resources:
`terraform apply/destroy -target=resourceName`
 - Caution: Dependencies are ignored!
- Remove resource from State:
`terraform state rm addressOfTheResource`
 - Resource is only removed from state but not destroyed
 - Use only as last resort!

DEMO: STATE MANIPULATION

TERRAFORM

USING TERRAFORM IN A TEAM

TERRAFORM

➤ Using Terraform as a team

- Till now only worked on one device with a local state
- In a team we need:
 - A distributed state
 - A way to prevent multiple state manipulations at the same time
 - Decide on a file structure
 - Integrate Terraform into the build chain

➤ Terraform Remote State

- Currently: Local state file for every team member
- Problems:
 - Possible outdated state file
 - Need to make sure no one modifies the state at the same time
- Solution: Remote state

“With remote state, Terraform writes the state data to a remote data store, which can then be shared between all members of a team.”

Supported state stores:

- Terraform Cloud
- HashiCorp Consul
- Amazon S3
- And more...

➤ Terraform Remote State

```
terraform {  
  backend "s3" {  
    bucket      = "myBucketName"  
    key         = "live/prodyna-aws-training/prod/mgmt/dns/terraform.tfstate"  
    region      = "eu-central-1"  
    dynamodb_table = "myDynamoDBTable"  
    encrypt     = true  
  }  
}
```

➤ Amazon S3 Remote State

- Amazon S3 for store the remote state
 - Encryption
 - Versioning
 - Access Control
 - High durability and availability
 - Serverless
- Amazon DynamoDB for state lock management
- Known issue! Need to export AWS credentials:
`export AWS_ACCESS_KEY_ID=XXXXXXXX`
`export AWS_SECRET_ACCESS_KEY=YYYYYYYY`

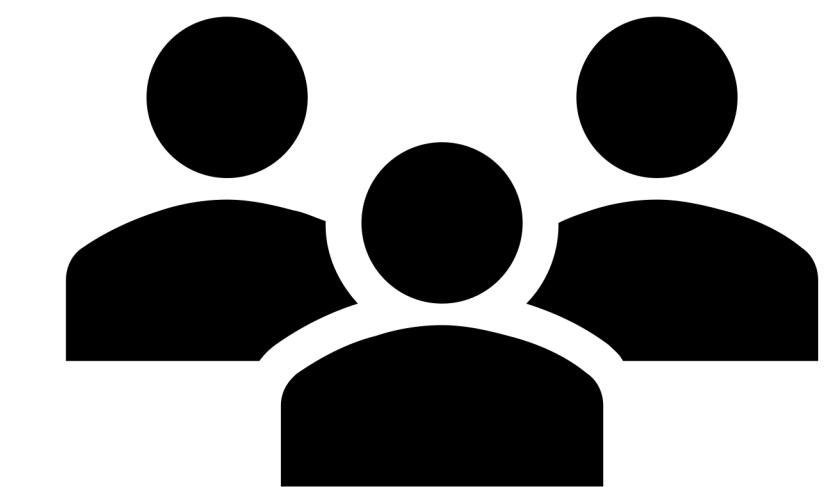
PRACTICE: MIGRATE TO REMOTE STATE

TERRAFORM

1. Create a remote backend with state lock management

- Create a S3 bucket with versioning and encryption enabled
- Create a Dynamo DB with

```
attribute {  
    name = "LockID"  
    type = "S"  
}  
  
hash_key = "LockID"
```

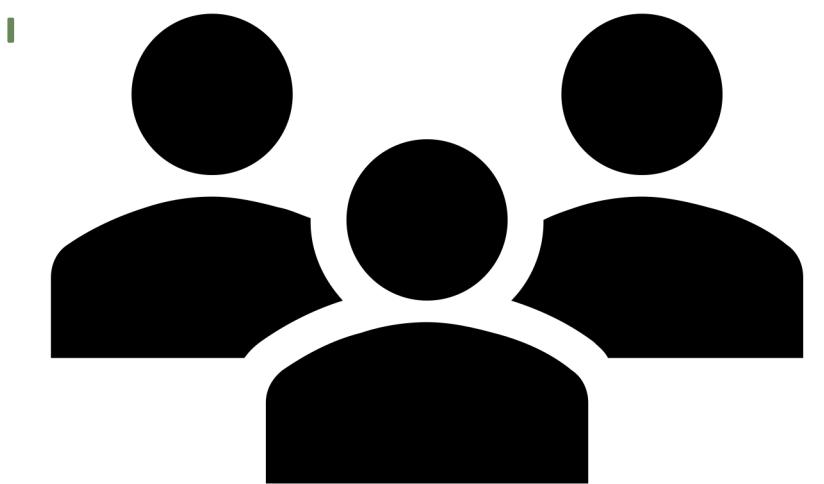


Time: 45 min

➤ Practice: Terraform in a Team

2. Migrate all local states to remote state. Example:

```
terraform {  
    backend "s3" {  
        bucket      = "myBucketName"  
        key         = "live/prodyna-aws-training/prod/mgmt/dns/terraform.tfstate"  
        region      = "eu-central-1"  
        dynamodb_table = "myDynamoDBTable"  
        encrypt     = true  
    }  
}
```



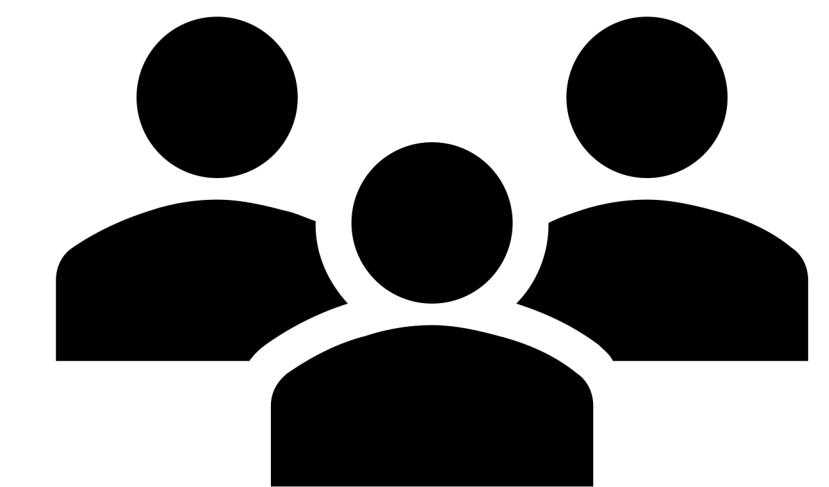
Time: 45 min

Integrate your team name and the stage „prod“ in your remote backend keys

➤ Practice: Terraform in a Team

3. Add version lock to prevent breaking changes. E.g.:

```
terraform {  
    required_version = "~> 0.12"  
    required_providers {  
        aws = "~> 2.47"  
    }  
}
```



Time: 45 min

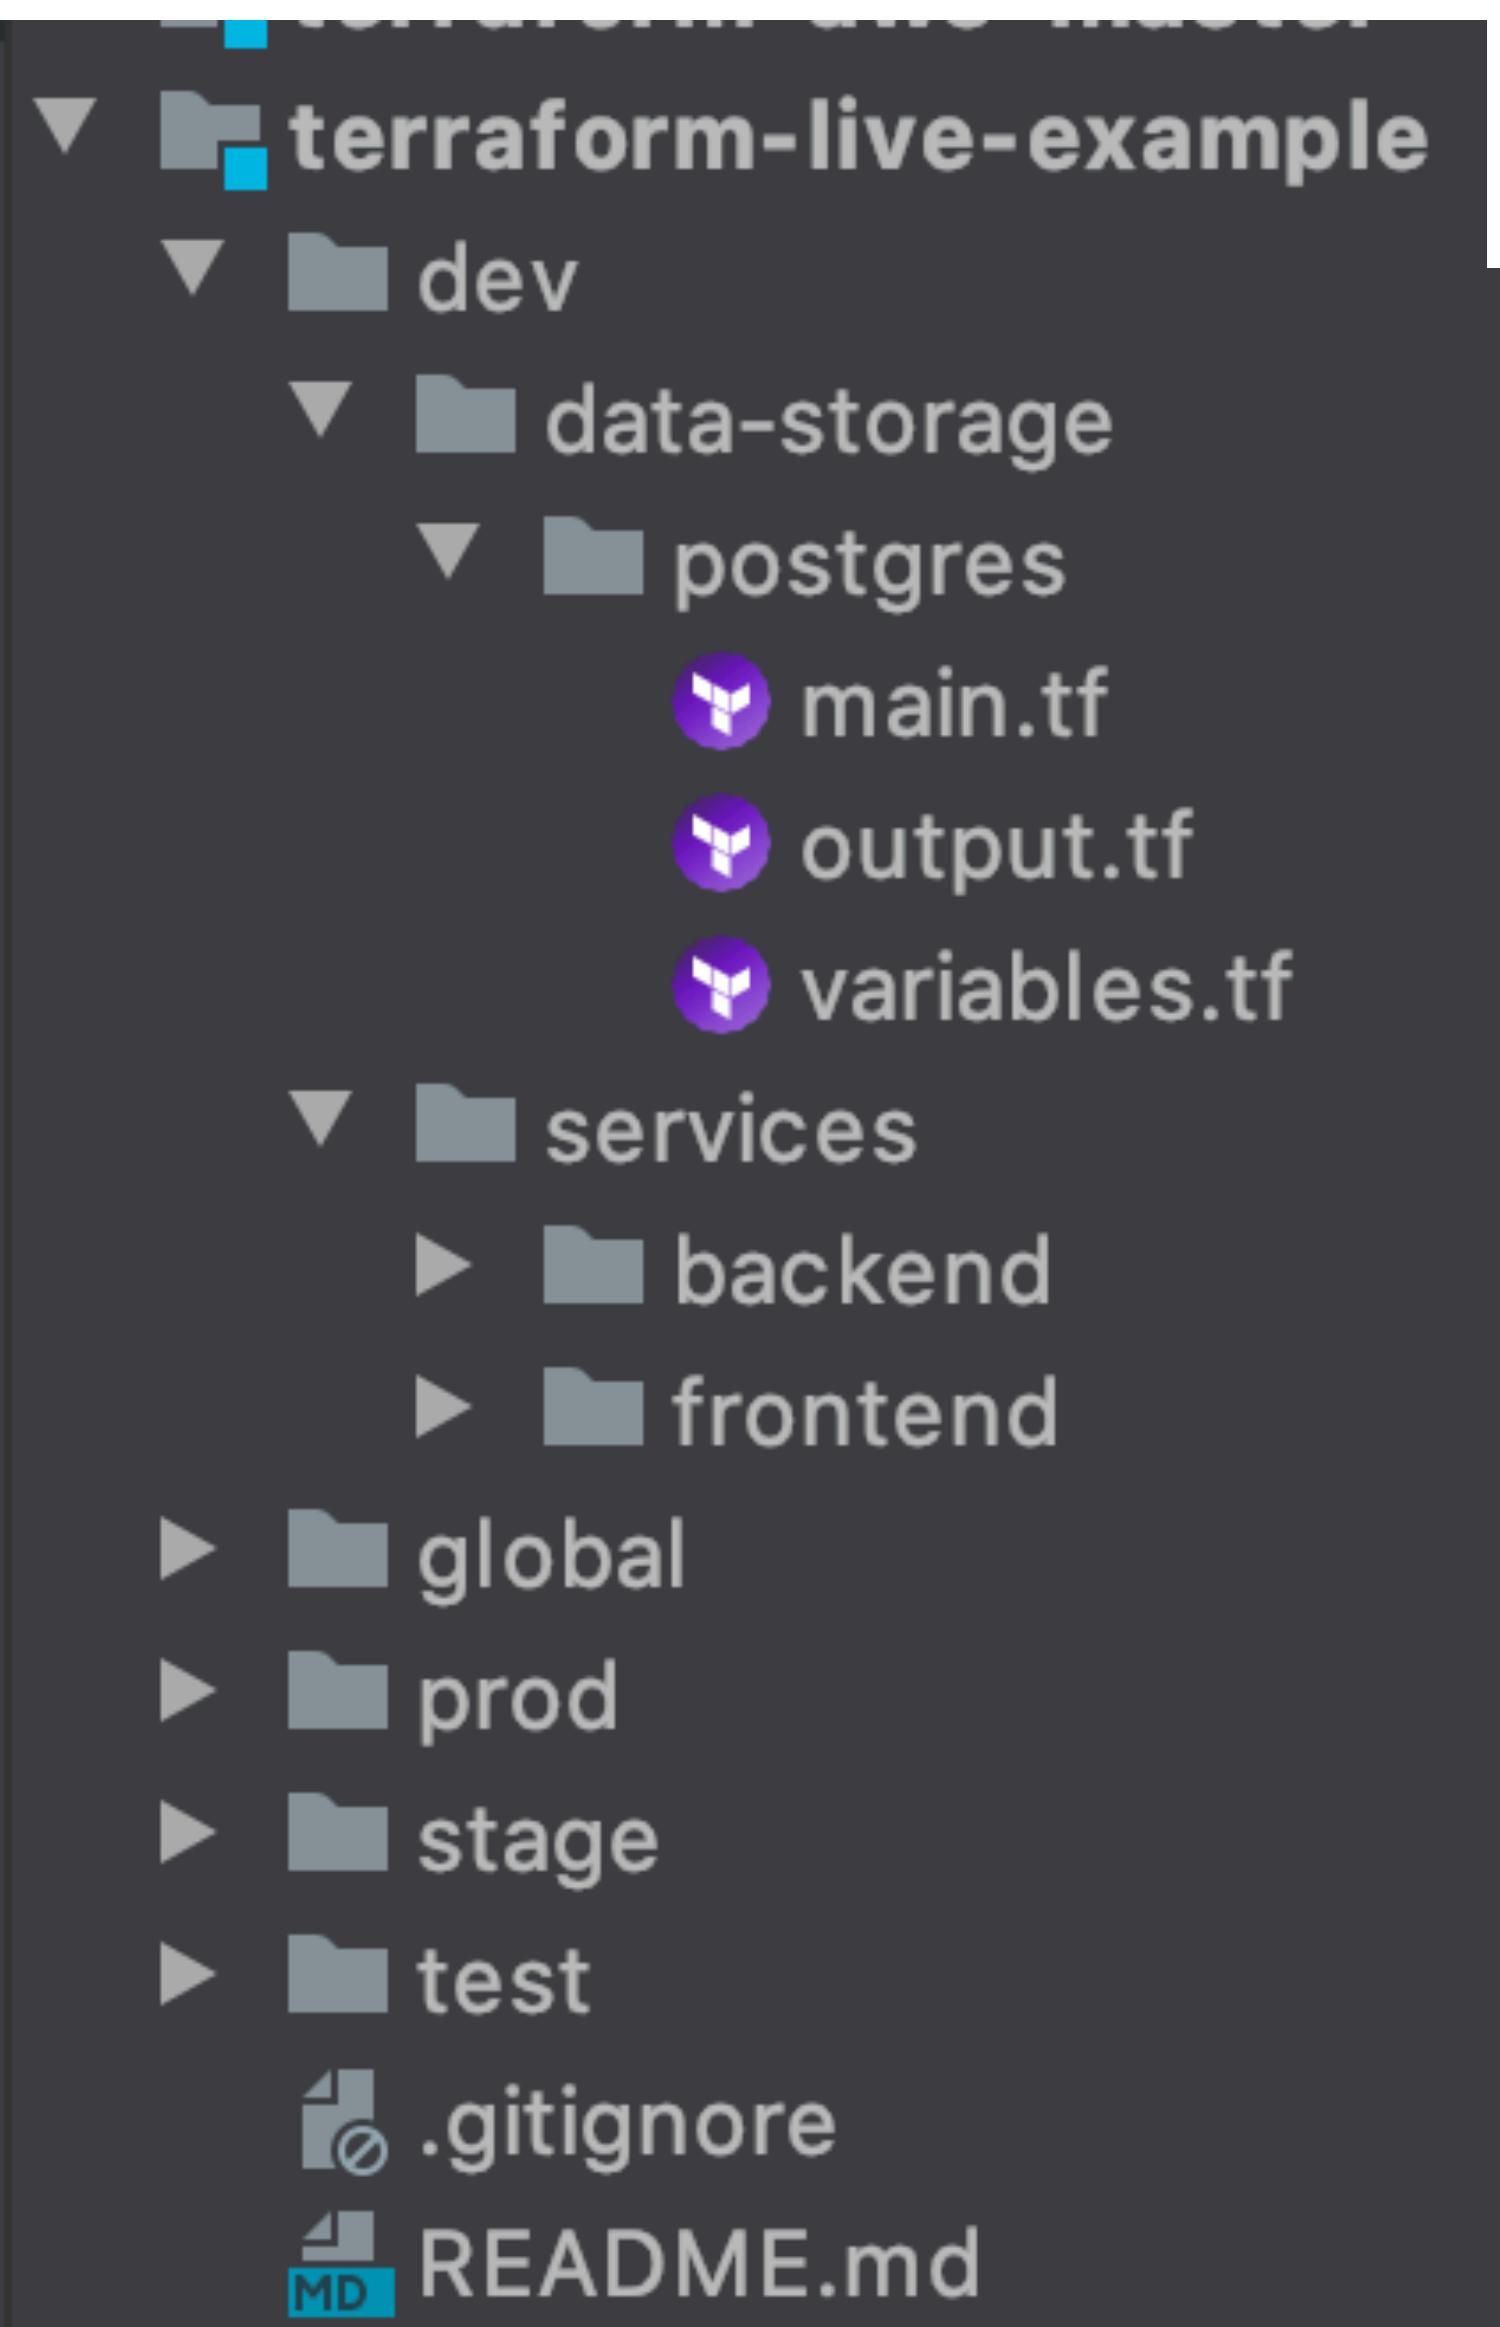
TERRAFORM FILE STRUCTURE

TERRAFORM

- Question: How to structure your Terraform files for multiple stages, projects and modules?
- Many answers, today only one

- Every project = Own git repository
- All modules in one git repository
 - Use tags for module versioning!
 - Lock module version with "?ref=myTag":

```
module "database" {  
    source = "git::https://mygitrepro.de?ref=1.1.17"
```
- Every stage (dev, test, prod, ...) = Own directory in repository (live repos)
- Each component (data-storage, services, ...) = Own directory in stage
- Use modules to avoid code duplication



terraform-modules

▼ data-stores

► managedPostgresDB

▼ mgmt

► clusterLogging

► clusterMonitoring

► sonarqube

► standardBillingAlert

► standardVPC

▼ security

▼ bastionHost

► aws.tf

► gcp.tf

► outputs.tf

► README.md

► variables.tf

► parameterDecrypter

► parameterStore

▼ services

► awsManagedKubernetesCluster

► gcpManagedKubernetesCluster

► terraformStateManager

► .gitignore

► README.md

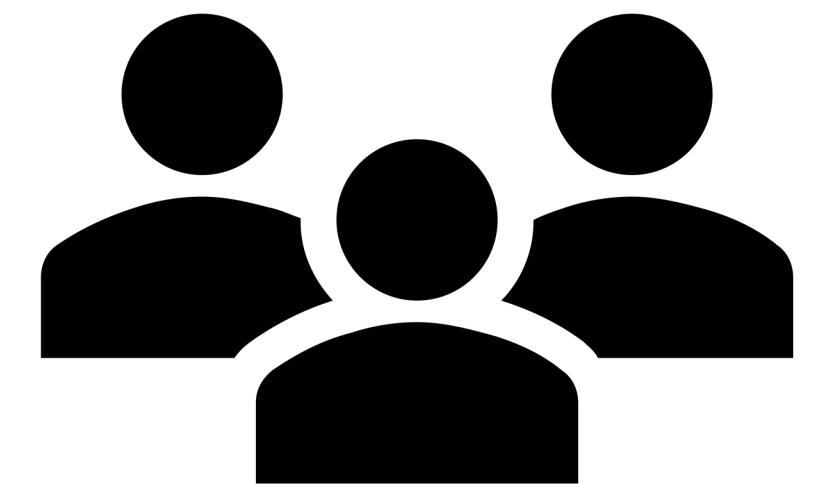
PRACTICE: TERRAFORM FILE STRUCTURE

TERRAFORM

➤ Practice: File Structure

Create a “live” and “modules” repository:

- Extract your modules to the new “modules” repository
- Apply the best practice file structure to your “live” repository



Time: 25 min

TERRAFORM IN YOUR BUILD CHAIN

TERRAFORM

➤ Integrate Terraform into your build chain

- Problem: Git branch from live repro. = different branch but same state file
- Solution: Terraform Workspaces
 - `terraform workspace show`
 - `terraform workspace new workspaceName`
 - `terraform workspace select workspaceName`
- Workspaces creates a new EMPTY state file
 - S3 path: `:env/xxx`
- You can now create and test your own infrastructure and destroy afterwards
- To avoid naming conflict add the workspace name to your resource names. E.g.:
`name = "${local.stageName}-${local.projectName}-${terraform.workspace}-s3-frontend"`

➤ Integrate Terraform into your build chain

- Question: How to integrate infrastructure changes?
- Answer:
 1. Create branch
 2. Create workspace
 3. Add changes
 4. Test changes
 5. Select default workspace
 6. Create Terraform plan and save output:
 - `terraform plan -out=myFeature.tfplan`
 - Save terminal output to file (`myFeature.tfplan.txt`)
 7. Create PR and add `*.tfplan` and `*.tfplan.txt`
 8. Merge



MISC

TERRAFORM

- You can import already created Cloud resources
- See provider docs for command
- E.g.: `terraform import aws_instance.web i-12345678`
- You need to have a resource with this name already in your Terraform file

- Execute action on local or remote machine
- Should be used as last resort – Use providers if possible
- Built-in provisioners:
 - Local-exec
 - Remote-exec
 - File
 - Puppet
 - ...

```
resource "aws_instance" "web" {  
...  
provisioner "local-exec" {  
  command = "echo The server's IP address is ${self.private_ip}"  
}  
}
```

- Use „self“ to access parent resource
- Use
`when = "destroy"`
to execute provisioner only on destroy

➤ Template Provider

„The template provider exposes data sources to use templates to generate strings for other Terraform resources or outputs.“

```
data "template_file" "init" {  
    template = file("${path.module}/init.tpl")  
    vars = {  
        consul_address = aws_instance.consul.private_ip  
    }  
}
```

- „file“ function loads a file
- `path.module` returns the full path of the current module

➤ Thank you



ANY QUESTIONS?